

Assignment 2 – Text part.

Tycho Canter Cremers

All code can be found in the weekly 2 folder.

Task I – CUDA programming

I.1

I.2

I.3

I timed the futhark sequential version and flat parallel version to compare them. On the standard suggested data set given in the comments the sequential version was just slightly faster than the parallel version. When I increased the dataset 10 fold the results below showed that in that case the parallel version is about twice as fast.

```
[mns267@a00332 SpMatVecMul-Futhark]$ futhark-dataset --i32-bounds=0:9999 -g [10000000]i32 --f32-bounds=-7.0:7.0 -g [10000000]f32 --i32-bounds=1000:1000 -g [10000]i32 --f32-bounds=-10.0:10.0 -g [10000]f32 | ./spMVmult-flat -r 10 -t /dev/stderr
7269
7232
7394
7293
7297
7276
7265
7346
7258
7307
```

seq

```
[mns267@a00332 SpMatVecMul-Futhark]$ futhark-dataset --i32-bounds=0:9999 -g [10000000]i32 --f32-bounds=-7.0:7.0 -g [10000000]f32 --i32-bounds=1000:1000 -g [10000]i32 --f32-bounds=-10.0:10.0 -g [10000]f32 | ./spMVmult-seq -r 10 -t /dev/stderr
15098
15053
15061
15049
15099
15094
15100
15033
15037
15097
```

Task II – Hardware

II.1

A.

I hope this is correct. I honestly can't find a good resource on where to actually put those NOOPS. I'm assuming you need one after every dependency hazard.

SEARCH:

LW R5, 0(R3)	/I1 Load item
SUB R6, R5, R2	/I2 compare with key
NOOP	
BNEZ R6, NOMATCH	/I3 check for match
NOOP	
ADDI R1, R1, #1	/I4 count matches

NOMATCH:

ADDI R3, R3, #4	/I5 next item
BNE R4, R3, SEARCH	/I6 continue until all items

B.

Again, I really hoping I'm doing this right.

MATCH

			c1	c2	c3	c4	c5	c6	c7	c8	c9	c10	c11	c12	c13	c14
SEARCH	L1	LW R5,0(R3)	F	D	EX	ME	WB									
	L2	SUB R6, R5, R2		F	D	D	D	EX	ME	WB						
	L3	BNEZ, R6, NOMATCH			F	F	F	D	D	D	EX	ME	WB			
	L4	ADDI R1, R1, #1					F	F	F	F	D	EX	ME	WB		
NOMATCH	L5	ADDI R3, R3, #4									F	D	EX	ME	WB	
	L6	BNE R4, R3, SEARCH										F	D	EX	ME	WB

**NO
MATCH**

			c1	c2	c3	c4	c5	c6	c7	c8	c9	c10	c11	c12	c13	c14
SEARCH	L1	LW R5,0(R3)	F	D	EX	ME	WB									
	L2	SUB R6, R5, R2		F	D	D	D	EX	ME	WB						
	L3	BNEZ, R6, NOMATCH			F	F	F	D	D	D	EX	ME	WB			
	L4	ADDI R1, R1, #1					F	F	F	F	NOP					
NOMATCH	L5	ADDI R3, R3, #4									F	D	EX	ME	WB	
	L6	BNE R4, R3, SEARCH										F	D	EX	ME	WB

L2 has to wait until the c5 WB to be able to access R5. Following, L3 has to wait for L2's writeback until he can access R6. This creates a huge delay.

So in both cases (MATCH and NO MATCH) it takes 14 clocks because you don't gain any time from stopping L4.

C.

Due to register forwarding L2 now only has to wait until c4 to access R5. Additionally, L3 has to wait one less cycle too, to access R6 at c7. This saves 2 clocks. So both cases can now do with 12 clocks.

D.

D MATCH			c1	c2	c3	c4	c5	c6	c7	c8	c9	c10	c11	c12	c13	c14
SEARCH	L1	LW R5,0(R3)	F	D	EX	ME	WB									
	L2	SUB R6, R5, R2		F	D	D	EX	ME	WB							
	L3	BNEZ, R6, NOMATCH			F	F	D	EX	ME	WB						
	L4	ADDI R1, R1, #1					F	D	EX	ME	WB					
NOMATCH	L5	ADDI R3, R3, #4						F	D	EX	ME	WB				
	L6	BNE R4, R3, SEARCH							F	D	EX	ME	WB			

D NO MATCH			c1	c2	c3	c4	c5	c6	c7	c8	c9	c10	c11	c12	c13	c14
SEARCH	L1	LW R5,0(R3)	F	D	EX	ME	WB									
	L2	SUB R6, R5, R2		F	D	D	EX	ME	WB							
	L3	BNEZ, R6, NOMATCH			F	F	D	EX	ME	WB						
	L4	ADDI R1, R1, #1					F	mop								
NOMATCH	L5	ADDI R3, R3, #4						F	D	EX	ME	WB				
	L6	BNE R4, R3, SEARCH							F	D	EX	ME	WB			

Due to full forwarding L3 gets R6 from the start of the EX phase in L2, saving another clock cycle.
Now both cases can do it in 11 clocks.

E.

Total number of clocks of 2 iterations = $11 * 2 = 22$ clocks.

If I did it correctly, unrolling the loop twice combined with register forwarding should remove all the dependency hazards for matches and not matches. So the total amount of clocks for a loop unrolled twice is 16:

E

MATCH

			c1	c2	c3	c4	c5	c6	c7	c8	c9	c10	c11	c12	c13	c14	c15	c16
SEARCH	L1	LW R5,0(R3)	F	D	EX	ME	WB											
	L2	LW R52,0(R32)		F	D	EX	ME	WB										
	L3	SUB R6, R5, R2			F	D	EX	ME	WB									
	L4	SUB R62, R52, R22				F	D	EX	ME	WB								
NOMATCH	L5	BNEZ, R6, NOMATCH					F	D	EX	ME	WB							
	L6	ADDI R1, R1, #2						F	D	EX	ME	WB						
	L7	ADDI R3, R3, #8							F	D	EX	ME	WB					
NOMATCH	L8	BNEZ, R62, NOMATCH							F	D	EX	ME	WB					
	L9	ADDI R12, R12, #2								F	D	EX	ME	WB				
	L10	ADDI R32, R32, #8									F	D	EX	ME	WB			
	L11	BNE R4, R3, SEARCH										F	D	EX	ME	WB		
	L12	BNE R42, R32, SEARCH											F	D	EX	ME	WB	

NO

MATCH

			c1	c2	c3	c4	c5	c6	c7	c8	c9	c10	c11	c12	c13	c14	c15	c16
SEARCH	L1	LW R5,0(R3)	F	D	EX	ME	WB											
	L2	LW R52,0(R32)		F	D	EX	ME	WB										
	L3	SUB R6, R5, R2			F	D	EX	ME	WB									
	L4	SUB R62, R52, R22				F	D	EX	ME	WB								
NOMATCH	L5	BNEZ, R6, NOMATCH					F	D	EX	ME	WB							
	L6	ADDI R1, R1, #2						F	nop									
	L7	ADDI R3, R3, #8							F	D	EX	ME	WB					
NOMATCH	L8	BNEZ, R62, NOMATCH							F	D	EX	ME	WB					
	L9	ADDI R12, R12, #2								F	nop							
	L10	ADDI R32, R32, #8									F	D	EX	ME	WB			
	L11	BNE R4, R3, SEARCH										F	D	EX	ME	WB		
	L12	BNE R42, R32, SEARCH											F	D	EX	ME	WB	

Which means that unrolling the loop would be usefull.

I don't think delayed branches would help because they didn't cause any delays in the first place.

Taske II.2

A.

ADDI R6,R0,#1

```
L.V V1,0(R1),R6    //Load slice of vector x
L.V V2,0(R2),R6    //Load slice of vector y
MUL.V V3,V1,V2     //Multiply x and y
ADD.V V4,V4,V3     // P += x*y
SUBBI R1,R1,#512    //jump to next slices,
SUBBI R2,R2,#512    //i.e., 8*64
SUBBI R3,R3,#64
BNEZ R3,LOOP
```

B.

From https://www.cs.indiana.edu/~achauhan/Teaching/B649/2010-Spring/StudentPresns/Vector_processors.pdf

Execution time =

$$T_n = \left[\frac{n}{MVL} \right] \times (T_{\text{loop}} + T_{\text{start}}) + n \times T_{\text{chime}}$$

I can't find anything about loop overhead or Tloop but in the examples they use 15. So I'll go with that.

N = 1024

MVL = 64

Tloop = 15

Tstart = 2*30+10+5 = 75

Tchime = 3 (I think)

So $T_n = 16 * (15 + 75) + 1024 * 3 = 4512$

C.

$(1024*1024) = 1048576$ dot products to be calculated.

So it's $1048576 * \text{answer to b.}$

$1048576 * 4512 = 4\,731\,174\,912$