

Weekly 3 – MPHP

Tycho Canter Cremers

Task 1

Task 2

A.

Why is the outer loop not parallel?

Because `accum` is declared outside the loop the same memory location is accessed and written to in different iterations possibly at the same time. So “`B[i,0] = accum;`” has a possibility of using the wrong value. To solve it you should privatize `accum` and declare it inside the loop.

```
1. for i from 0 to N-1 // outer loop
2.   accum;
3.   accum = A[i,0] * A[i,0];
4.   B[i,0] = accum;
5. for j from 1 to 63 // inner loop
6.   tmpA = A[i, j];
7.   accum = sqrt(accum) + tmpA*tmpA;
8.   B[i,j] = accum;
```

Is the inner loop parallel?

Again, no. Because `accum` is, once again declared outside the loop. It could be written too, non sequentially, in another iteration resulting in the wrong value at `B[i,j]`. Additionally, `sqrt` is not an associative operation, so can't be used in parallel in this case.

Can the inner loop be re-written as a composition of parallel operators?

No, because taking the square root of `accum` is not associative.

If line 6 is re-written as `accum = accum + tmpA*tmpA`, would it be possible now to rewrite the inner loop as a composition of bulk operators?

Now you can. It's basically a map to calculate `tmpA*tmpA` and then a scan.

D.

When you use the make file to run both of the original Program and the transformed program I get this output. The Transformed program is almost 10 times faster than the original, non-transposed one.

```
[mns267@a00332 Task1-2]$ ./matrixMult_2
```

```
Original Program on GPU runs in: 46255 microsecs
```

```
GPU PROGRAM  VALID!
```

```
[mns267@a00332 Task1-2]$ ./matrixMult_3
```

```
Transformed Program on GPU runs in: 5154 microsecs
```

```
GPU PROGRAM  VALID!
```

Task 3

d.

```
./matrixMult
```

```
Sequential Naive version runs in: 5393490 microsecs
```

```
GPU version runs in: 5729 microsecs
```

```
Performance= 374.84 GFlop/s, Time= 5729.000 microsec 64 64
```

```
VALID RESULT!
```

Above is the output for the code ran with the tiled kernel. The tiled parallel version is almost 1000 times faster than the sequential version.

```
./matrixMult
```

```
Sequential Naive version runs in: 5409844 microsecs
```

```
GPU version runs in: 17280 microsecs
```

```
Performance= 124.28 GFlop/s, Time= 17280.000 microsec 64 64
```

```
VALID RESULT!
```

Running the code with the normal kernel runs about a little more than 3 times slower than the tiled kernel. But still a lot faster than the sequential version.